

Loopy belief propagation

CS B553
Spring 2013


Announcements

- A2 due tonight
 - Reminder of late policy: 10% off for up to 48 hours late
- A3 coming soon!

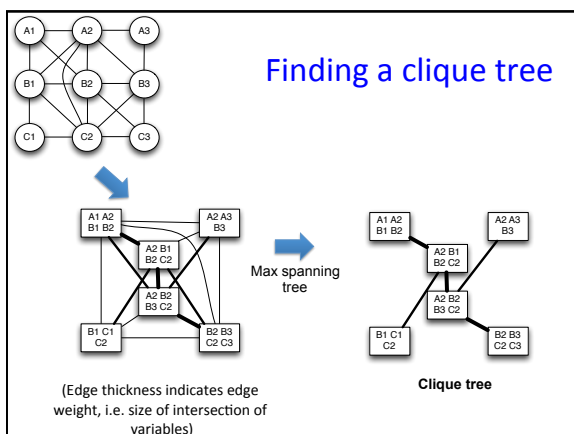
Constructing a clique tree

- One way to find a clique tree is to choose a variable elimination ordering and “run” VE
- Another approach is to use a graph construction
 - If necessary, **moralize** to produce an undirected graph G.
 - **Triangulate** G to produce a chordal graph H. (Would like one with minimum clique size, but this is NP hard.)
 - Find **maximal cliques** in H. (Not NP hard for chordal graphs.)
 - **Construct graph** with nodes corresponding to max cliques in H, edges weighted according to degree of overlap. i.e. edge between C1 and C2 has weight $|C1 \cap C2|$
 - Find a **max spanning tree** on this graph to yield a clique tree.

Clique trees for grid graphs



Finding a clique tree



Tree width

- The **tree width** of a graph G is equal to $m-1$,
 - Where m is the size of the largest clique in the triangulated (chordal) version of G
- The worst-case running time of exact inference on a Markov or Bayes network is exponential in the tree width of the moralized graph.
 - E.g., the tree width of an $n \times n$ grid graph is n , so inference takes $O(s^n)$ time, where s is # of possible values of each random variable

Making inference tractable

- Three strategies for efficient inference in practice:
 1. Use a graphical model with **low tree width**.
 2. Make some additional assumptions about the problem.
(Efficient algorithms exist for some special cases, e.g. when the clique potentials have a specific form.)
 3. Settle for an approximate inference algorithm.


Making inference tractable

- In practice, making inference tractable is a key challenge in applying graphical models to applications
- Typically, the options are:
 - Exact inference with **arbitrary potentials** on a graphical model, but with a **simplified structure**
 - Exact inference on a graphical model with **arbitrary structure**, but **restricted potentials**
 - Graphical model with **arbitrary structure and arbitrary potentials**, but settle for **approximate inference**


Approximate inference

- So far, we've covered *exact* inference algorithms
 - I.e. algorithms that can compute marginal probability distributions exactly, with no error
 - They are fast for some graphical models, like trees, but are prohibitively expensive for most applications
- But is exact inference worth it?
 - Our models aren't perfect anyway...
- Much more efficient *approximate* inference algorithms exist
 - If we're willing to settle for inexact answers

Consider a simple cycle:




- A clique tree for this cycle is:




- Recall that a clique tree is a special *cluster graph* (each node is a clique of the original graph, edges connect subset of nodes sharing a common variable)

- Here's another cluster graph for the above cycle:



- It's not a clique tree
- It's a *loopy* graph

Back to sum-product...



- Recall that sum-product involves sending messages to neighbors,
- $$\delta_{i \rightarrow j}(\mathbf{S}_{i,j}) = \sum_{\mathbf{C}_i - \mathbf{C}_j} \psi_i(\mathbf{C}_i) \prod_{k \in \mathcal{N}(i) - \{j\}} \delta_{k \rightarrow i}(\mathbf{S}_{k,i})$$

- We derived this for use on a clique tree, but the definition of messages (above) works on any graph!
 - I.e., there's nothing that prevents us from running Sum-Product Belief Propagation on a loopy cluster graph
 - There's just no theoretical justification for doing this...

Sum product

$$\delta_{i \rightarrow j}(\mathbf{S}_{i,j}) = \sum_{\mathbf{C}_i - \mathbf{C}_j} \psi_i(\mathbf{C}_i) \prod_{k \in \mathcal{N}(i) - \{j\}} \delta_{k \rightarrow i}(\mathbf{S}_{k,i})$$


BP on a clique tree

- Messages start at leaves, propagate up to root, then propagate down to leaves



BP on loopy cluster graph

- Nodes send messages to their neighbors, based on their clique potential and messages from neighbors



Loopy Belief Propagation

- Construct a cluster graph from the Markov network
- Each node sends messages to its neighbor,

$$\delta_{i \rightarrow j}(\mathbf{S}_{i,j}) = \sum_{\mathbf{C}_i - \mathbf{C}_j} \psi_i(\mathbf{C}_i) \prod_{k \in \mathcal{N}(i) - \{j\}} \delta_{k \rightarrow i}(\mathbf{S}_{k,i})$$

- Problems with this?
- How to begin?
 - Initialize by having every node send a “fake” initial message, consisting of just a uniform distribution
- How to end?
 - Keep running iterations of BP until convergence
 - Unfortunately, convergence might not happen...

Historical note: Turbo codes

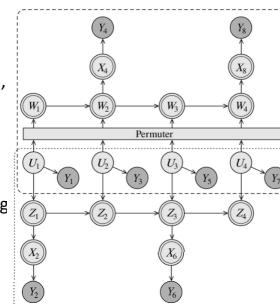
- Loopy Belief Propagation (LBP) has been known for ~25 years, but people assumed it was not useful
- In the 1990s, a seemingly unrelated discovery sparked renewed interest in LBP
 - Paper on a new error-correcting transmission algorithm, for noisy links (e.g. wireless networks)
 - Breakthrough in communications technology; modern wireless phones still use this (or a related) technique

Noisy link model

- Suppose we want to send bits U_1, U_2, U_3, U_4 across a noisy communications link
 - Need some redundancy to detect and correct errors
 - So, we encode these bits as some new bits $X_1 - X_8$
 - The observer receives some bits $Y_1 - Y_8$, which may be a corrupted copy of $X_1 - X_8$
- From a probabilistic inference standpoint, receiver wants to compute distribution over \mathbf{X} given \mathbf{Y}

Turbo codes

- Berrou (1993)
- Solve for $U_1 - U_4$, given $Y_1 - Y_8$
- Hack: First solve for Y_1, Y_3, Y_5, Y_7 .
- Hold these constant, then solve for Y_2, Y_4, Y_6, Y_8 .
- Then repeat, iteratively; i.e. “Turbocharge” estimates using iterative feedback loop
- Unbeknownst to them, they had (re-)discovered Loopy Belief Propagation!



Loopy BP properties

- Recall that clique trees exhibited the *running intersection property*:
 - If variable X is part of node C and node D , then every node along the path from C to D also contains X
- Loopy cluster graphs won't satisfy this condition
- Instead, we'll require the following weaker property:
 - If variable X is part of node C and node D , then there exists exactly one path between C and D such that the scope of the messages along the path include X
 - Intuitively, helps to prevent endless propagation cycles

Loopy BP properties

- Loopy BP can be **much** faster than exact inference
 - Running time?
- Loopy BP is not guaranteed to converge to a solution
- Even if it does converge, it's not guaranteed to converge to the correct solution
 - However, in practice it seems to converge to a reasonably good solution for a variety of Markov net problems
- Rather little is known about its theoretical properties
 - Including when it will work well and when it won't
 - We'll see (later) one explanation for why it works

Loopy BP on grid graphs

- In the special case of grid graphs, we can run BP on the Markov net directly, instead of the cluster graph
- This is possible because we can create a cluster graph with a very simple structure
- The computation is exactly the same, but it's more convenient to understand and implement this way

Belief propagation

- Messages are passed between the variables

– Message from node i to j at iteration t is,

$$m_{i \rightarrow j}^{(t)}(X_j) = \sum_{X_i} \phi(X_i, X_j) \prod_{k \in \mathcal{N}(X_i) - \{X_j\}} m_{k \rightarrow i}^{(t-1)}(X_i)$$

– Intuitively (and anthropomorphically) a message from me to you says, for each of your states, “If you choose state X_j , here’s how happy my neighbors and I would be.”

