

Graph cuts


CS B553
Spring 2013

Faster MAP inference?

- We've now seen two algorithms for MAP inference:
 - Variable elimination: Exact, but potentially very slow
 - Loopy Max-product BP: Fast, but approximate
- It turns out that in some cases, MAP problems are easier than Marginal inference problems
 - One interesting case: With binary random variables, and potential functions that satisfy (relatively weak) restrictions, exact inference on a pairwise Markov network is efficient

Consider a simpler problem...

- Suppose we want to find a bright object against a dark background
 - But some of the pixel values are slightly wrong
 - So want to estimate a 0 or 1 label for each pixel



A slightly more interesting problem...


- Foreground vs background segmentation
 - We want to label every pixel of an image with a 0 or a 1, indicating whether it's a background or foreground pixel



Adapted from N. Snavely's slide


A slightly more interesting problem...

- One approach: Use some human input
 - Use sketches out a few strokes to indicate foreground and background, then we try to classify rest of pixels.




Adapted from N. Snavely's slide

Solving with an MRF




Observed pixel data



Unobservable binary labels

Data cost



D(fg, Y)



D(bg, Y)

Adapted from N. Snavely's slide

Solving with an MRF

- So, we want to solve a problem of the form:


$$X^* = \arg \min_X \sum_i D_i(X_i, Y_i) + \sum_{(i,j) \in \mathcal{E}} V(X_i, X_j)$$

- Y variables are given
- X variables are binary-valued
- D cost functions have any form
- V cost functions have the form:

$$V(X_i, X_j) = \begin{cases} 0 & \text{if } X_i = X_j \\ k & \text{otherwise} \end{cases}$$



Observed pixel data



Unobservable binary labels


Network flow can help

- The minimization problem with 2 labels can be solved exactly using network flow
 - Construction probably due to [Hammer et al. 65]
 - First applied to images by [Greig et al. 86]
- Classical Computer Science problem reduction
 - Turn a new problem into a problem we can solve!

9

Adapted from R. Zabih's slide

Maximum flow problem




A graph with two terminals

- Max flow problem:
 - Each edge is a "pipe"
 - Find the largest flow F of "water" that can be sent from the "source" to the "sink" along the pipes
 - Source output = sink input = flow value
 - Edge weights give the pipe's capacity

10

Adapted from R. Zabih's slide

Minimum cut problem




A graph with two terminals

- Min cut problem:
 - Find the cheapest way to cut the edges so that the "source" is separated from the "sink"
 - Cut edges going from source side to sink side
 - Edge weights now represent cutting "costs"

11

Adapted from R. Zabih's slide

Max flow/Min cut theorem




A graph with two terminals

- Max Flow = Min Cut:
 - Proof sketch: value of a flow is value over any cut
 - Maximum flow saturates the edges along the minimum cut
 - Ford and Fulkerson, 1962
 - Problem reduction!
- Ford and Fulkerson gave first polynomial time algorithm for globally optimal solution

12

Adapted from R. Zabih's slide

"Augmenting Path" algorithms




A graph with two terminals

13

Adapted from R. Zabih's slide

- Find a path from S to T along non-saturated edges
- Increase flow along this path until some edge saturates

"Augmenting Path" algorithms




A graph with two terminals

14

Adapted from R. Zabih's slide

- Find a path from S to T along non-saturated edges
- Increase flow along this path until some edge saturates
- Find next path...
- Increase flow...

"Augmenting Path" algorithms



A graph with two terminals

15

Adapted from R. Zabih's slide

- Find a path from S to T along non-saturated edges
- Increase flow along this path until some edge saturates

Iterate until all paths from S to T have at least one saturated edge

Min flow algorithms

- Ford-Fulkerson (1962) is the classic algorithm
 - Takes time $O(|E| f)$, where f is the maximum flow
 - May not converge in some cases
- Edmonds-Karp (1972) gave an improved version
 - Same as F-F, but the augmented path is always the shortest with available capacity. Can be found using breadth-first search.
 - Takes time $O(|V| |E|^2)$

Adapted from R. Zabih's slide


Back to binary MRFs...

- We want to solve a problem of the form:

$$X^* = \arg \min_X \sum_i D_i(X_i, Y_i) + \sum_{(i,j) \in \mathcal{E}} V(X_i, X_j)$$

- Y variables are given
- X variables are binary-valued
- D cost functions have any form
- V cost functions have the form:


$$V(X_i, X_j) = \begin{cases} 0 & \text{if } X_i = X_j \\ k & \text{otherwise} \end{cases}$$



Unobservable binary labels

Basic graph cut construction

- One non-terminal vertex per pixel
 - Each pixel connects directly to s, t, and to its neighbors
 - Edge to t has weight $D_p(0)$, edge to s has weight $D_p(1)$
 - Edge (p,q) has weight $V_{pq}(0,1)$
- Cost of cut is the cost of the entire labeling
 - Pixel p labeled 1 if connected to t, labeled 0 if connected to s



18

Adapted from R. Zabih's slide

$$E(x_1, \dots, x_n) = \sum_p D_p(x_p) + \sum_{p,q} V_{p,q}(x_p, x_q)$$